

Quantium2

December 30, 2023

1 Importing Libraries

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
from scipy.stats import ttest_ind
from scipy.stats import pearsonr
```

2 Loading and Exploring Dataset

```
[2]: df = pd.read_csv("C:/Users/Asus/Desktop/Forage/QUANTUM DA VIRTUAL INTERNSHIP/
↳QVI_data.csv")

print(df.head()) # Display the first few rows of the DataFrame

print(df.describe()) # Display basic statistics of the data

print(df.info()) # Check data types and missing values
```

	LYLTY_CARD_NBR	DATE	STORE_NBR	TXN_ID	PROD_NBR	\
0	1000	10/17/2018	1	1	5	
1	1002	9/16/2018	1	2	58	
2	1003	3/7/2019	1	3	52	
3	1003	3/8/2019	1	4	106	
4	1004	11/2/2018	1	5	96	

	PROD_NAME	PROD_QTY	TOT_SALES	PACK_SIZE	\
0	Natural Chip Compny SeaSalt175g	2	6.0	175	
1	Red Rock Deli Chikn&Garlic Aioli 150g	1	2.7	150	
2	Grain Waves Sour Cream&Chives 210G	1	3.6	210	
3	Natural ChipCo Hony Soy Chckn175g	1	3.0	175	
4	WW Original Stacked Chips 160g	1	1.9	160	

	BRAND	LIFESTAGE	PREMIUM_CUSTOMER
0	NATURAL	YOUNG SINGLES/COUPLES	Premium

```

1      RRD  YOUNG SINGLES/COUPLES      Mainstream
2      GRNWVES      YOUNG FAMILIES      Budget
3      NATURAL      YOUNG FAMILIES      Budget
4  WOOLWORTHS  OLDER SINGLES/COUPLES      Mainstream
      LYLTY_CARD_NBR      STORE_NBR      TXN_ID      PROD_NBR  \
count      2.648340e+05  264834.000000  2.648340e+05  264834.000000
mean      1.355488e+05      135.079423  1.351576e+05      56.583554
std      8.057990e+04      76.784063  7.813292e+04      32.826444
min      1.000000e+03      1.000000  1.000000e+00      1.000000
25%      7.002100e+04      70.000000  6.760050e+04      28.000000
50%      1.303570e+05      130.000000  1.351365e+05      56.000000
75%      2.030940e+05      203.000000  2.026998e+05      85.000000
max      2.373711e+06      272.000000  2.415841e+06      114.000000

```

```

      PROD_QTY      TOT_SALES      PACK_SIZE
count  264834.000000  264834.000000  264834.000000
mean      1.905813      7.299346      182.425512
std      0.343436      2.527241      64.325148
min      1.000000      1.500000      70.000000
25%      2.000000      5.400000      150.000000
50%      2.000000      7.400000      170.000000
75%      2.000000      9.200000      175.000000
max      5.000000      29.500000      380.000000

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 264834 entries, 0 to 264833
```

```
Data columns (total 12 columns):
```

```

#      Column      Non-Null Count  Dtype
---  -
0  LYLTY_CARD_NBR  264834 non-null  int64
1  DATE            264834 non-null  object
2  STORE_NBR       264834 non-null  int64
3  TXN_ID          264834 non-null  int64
4  PROD_NBR        264834 non-null  int64
5  PROD_NAME       264834 non-null  object
6  PROD_QTY        264834 non-null  int64
7  TOT_SALES       264834 non-null  float64
8  PACK_SIZE       264834 non-null  int64
9  BRAND           264834 non-null  object
10 LIFESTAGE       264834 non-null  object
11 PREMIUM_CUSTOMER 264834 non-null  object

```

```
dtypes: float64(1), int64(6), object(5)
```

```
memory usage: 24.2+ MB
```

```
None
```

3 Identifying and Handling Outliers

```
[3]: ## Visualize the distribution of the variable of interest (e.g., 'spend')
sns.histplot(df['PROD_QTY'], kde=True)
plt.title('Distribution of PROD_QTY')
plt.show()

# Calculate Z-scores
z_scores = np.abs(stats.zscore(df['PROD_QTY']))

# Define a threshold for identifying outliers (e.g., Z-score > 3)
outliers = (z_scores > 3)

# Identify and print the outliers
outlier_values = df['PROD_QTY'][outliers]
print("Outlier values:")
print(outlier_values)

# Remove outliers from the dataset
df_no_outliers = df[~outliers]

# Display the box plot after removing outliers for comparison
sns.histplot(x=df_no_outliers['PROD_QTY'])
plt.show()

## Visualize the distribution of the variable of interest (e.g., 'spend')
sns.histplot(df['TOT_SALES'], kde=True)
plt.title('Distribution of TOT_SALES')
plt.show()

# Calculate Z-scores
z_scores = np.abs(stats.zscore(df['PROD_QTY']))

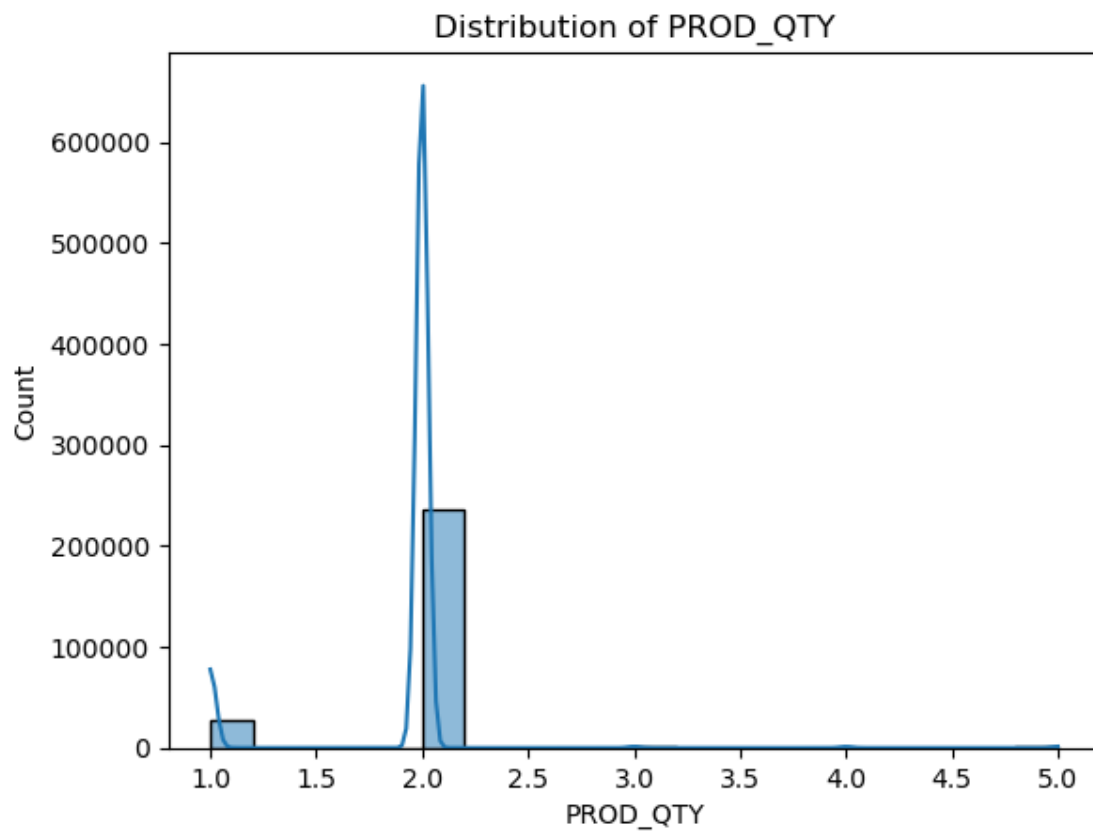
# Define a threshold for identifying outliers (e.g., Z-score > 3)
outliers = (z_scores > 3)

# Identify and print the outliers
outlier_values = df['TOT_SALES'][outliers]
print("Outlier values:")
print(outlier_values)

# Remove outliers from the dataset
df_no_outliers = df[~outliers]

# Display the box plot after removing outliers for comparison
```


```
sns.histplot(x=df_no_outliers['TOT_SALES'])
plt.show()
```

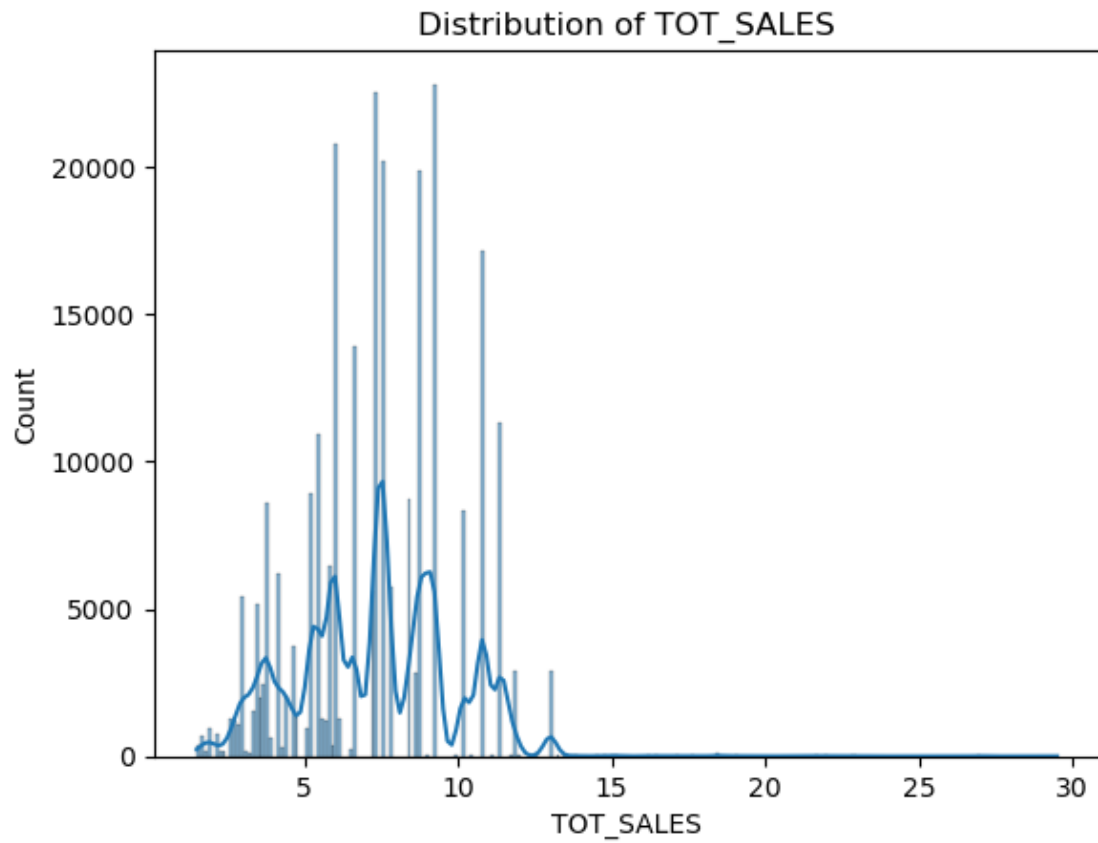


Outlier values:

315	4
332	3
523	4
952	5
1012	3
..	
263285	3
263431	3
263668	3
263830	4
264269	4

Name: PROD_QTY, Length: 1277, dtype: int64






Outlier values:

315	7.6
332	6.3
523	10.4
952	15.0
1012	13.8

...

263285	9.0
263431	8.1
263668	11.7
263830	10.8
264269	21.6

Name: TOT_SALES, Length: 1277, dtype: float64



4 Define Metrics

```
[4]: # Total number of customers
total_customers = df['LYLTY_CARD_NBR'].nunique()
print(f"Total number of customers: {total_customers}")
print()

# Average number of transactions per customer
average_transactions_per_customer = df.groupby('LYLTY_CARD_NBR')['TXN_ID'].
    ↪nunique().mean()
print(f"Average number of transactions per customer: ␣
    ↪{average_transactions_per_customer:.2f}")
print()

# Total sales revenue
total_sales_revenue = df['TOT_SALES'].sum()
print(f"Total sales revenue: {total_sales_revenue:.2f}")
print()

# Total Spend per Customer
```

```

total_spend_per_customer = df.groupby('LYLTY_CARD_NBR')['TOT_SALES'].sum()
df['total_spend_per_customer'] = total_spend_per_customer
print("Total Spend per Customer:")
print(total_spend_per_customer)
print()

# Average Spend per Customer
average_spend_per_customer = df.groupby('LYLTY_CARD_NBR')['TOT_SALES'].mean()
df['average_spend_per_customer'] = average_spend_per_customer
print("Average Spend per Customer:")
print(average_spend_per_customer)
print()

# Frequency of Purchase
purchase_frequency = df.groupby('LYLTY_CARD_NBR').size()
df['purchase_frequency'] = purchase_frequency
print("Frequency of Purchase:")
print(purchase_frequency)
print()

# Average Pack Size per Customer
average_pack_size_per_customer = df.groupby('LYLTY_CARD_NBR')['PACK_SIZE'].
    .mean()
df['average_pack_size_per_customer'] = average_pack_size_per_customer
print("Average Pack Size per Customer:")
print(average_pack_size_per_customer)
print()

```

Total number of customers: 72636

Average number of transactions per customer: 3.62

Total sales revenue: 1933115.00

Total Spend per Customer:

LYLTY_CARD_NBR

1000	6.0
1002	2.7
1003	6.6
1004	1.9
1005	2.8

...

2370651	13.0
2370701	7.2
2370751	9.2
2370961	18.6
2373711	11.4

Name: TOT_SALES, Length: 72636, dtype: float64

Average Spend per Customer:

LYLTY_CARD_NBR

1000	6.0
1002	2.7
1003	3.3
1004	1.9
1005	2.8

...

2370651	13.0
2370701	7.2
2370751	9.2
2370961	9.3
2373711	11.4

Name: TOT_SALES, Length: 72636, dtype: float64

Frequency of Purchase:

LYLTY_CARD_NBR

1000	1
1002	1
1003	2
1004	1
1005	1

..

2370651	1
2370701	1
2370751	1
2370961	2
2373711	1

Length: 72636, dtype: int64

Average Pack Size per Customer:

LYLTY_CARD_NBR

1000	175.0
1002	150.0
1003	192.5
1004	160.0
1005	165.0

...

2370651	380.0
2370701	210.0
2370751	150.0
2370961	232.5
2373711	330.0

Name: PACK_SIZE, Length: 72636, dtype: float64

5 Function to calculate the correlation or magnitude distance

```
[9]: def calculate_distance(trial_store, control_store, measure="correlation"):
    trial_data = df[df['STORE_NBR'] == trial_store]
    control_data = df[df['STORE_NBR'] == control_store]

    if measure == "correlation":
        return np.corrcoef(trial_data['TOT_SALES'],
        ↪control_data['TOT_SALES'])[0, 1]
    elif measure == "magnitude_distance":
        observed_distance = np.sum(np.abs(trial_data['TOT_SALES'] -
        ↪control_data['TOT_SALES']))
        min_distance = np.sum(np.abs(df['TOT_SALES'] -
        ↪control_data['TOT_SALES']))
        max_distance = np.sum(np.abs(df['TOT_SALES'] - df['TOT_SALES']))

        return 1 - (observed_distance - min_distance) / (max_distance -
        ↪min_distance)
```

6 Function to find control stores for a given trial store

```
[8]: def find_control_stores(trial_store, measure="correlation"):
    control_stores = set(df['STORE_NBR'].unique()) - {trial_store}

    distances = [calculate_distance(trial_store, control_store, measure) for
    ↪control_store in control_stores]

    selected_control_store = control_stores[np.argmax(distances)]

    return selected_control_store
```

7 Function to calculate a metric for comparing control stores to trial stores

```
[7]: # Define trial and control stores
trial_stores = [77, 86, 88]
control_stores = [store for store in df['STORE_NBR'].unique() if store not in
    ↪trial_stores]

# Function to calculate similarity metric
def calculate_similarity_metric(trial_store, control_store, metric='pearson'):
    trial_data = df[df['STORE_NBR'] == trial_store].select_dtypes(include=np.
    ↪number)
    control_data = df[df['STORE_NBR'] == control_store].
    ↪select_dtypes(include=np.number)
```

```

    if metric == 'pearson':
        similarity_metric = trial_data.corrwith(control_data, axis=0,
↪numeric_only=True)
    elif metric == 'magnitude_distance':
        # Implement your method for magnitude distance
        pass
    else:
        raise ValueError("Invalid metric. Choose 'pearson' or
↪'magnitude_distance'")

    return similarity_metric

# Function to select control stores
def select_control_stores(trial_store, control_stores, metric='pearson'):
    similarity_scores = []

    for control_store in control_stores:
        similarity_metric = calculate_similarity_metric(trial_store,
↪control_store, metric)
        avg_similarity_score = np.mean(similarity_metric)
        similarity_scores.append((control_store, avg_similarity_score))

    sorted_control_stores = sorted(similarity_scores, key=lambda x: x[1],
↪reverse=True)
    return sorted_control_stores[0][0]

# Function to perform t-test and analyze impact
def analyze_trial_effect(trial_store, control_store):
    trial_data = df[df['STORE_NBR'] == trial_store]
    control_data = df[df['STORE_NBR'] == control_store]

    t_stat, p_value = ttest_ind(trial_data['TOT_SALES'],
↪control_data['TOT_SALES'])

    if p_value < 0.05:
        print(f"Trial Store {trial_store} has a significant impact on total
↪sales compared to Control Store {control_store}.")
    else:
        print(f"No significant impact on total sales observed between Trial
↪Store {trial_store} and Control Store {control_store}.")

# Apply the functions for each trial store
for trial_store in trial_stores:
    selected_control_store = select_control_stores(trial_store, control_stores,
↪metric='pearson')

```

```
analyze_trial_effect(trial_store, selected_control_store)
```

Trial Store 77 has a significant impact on total sales compared to Control Store 1.

Trial Store 86 has a significant impact on total sales compared to Control Store 1.

Trial Store 88 has a significant impact on total sales compared to Control Store 1.